

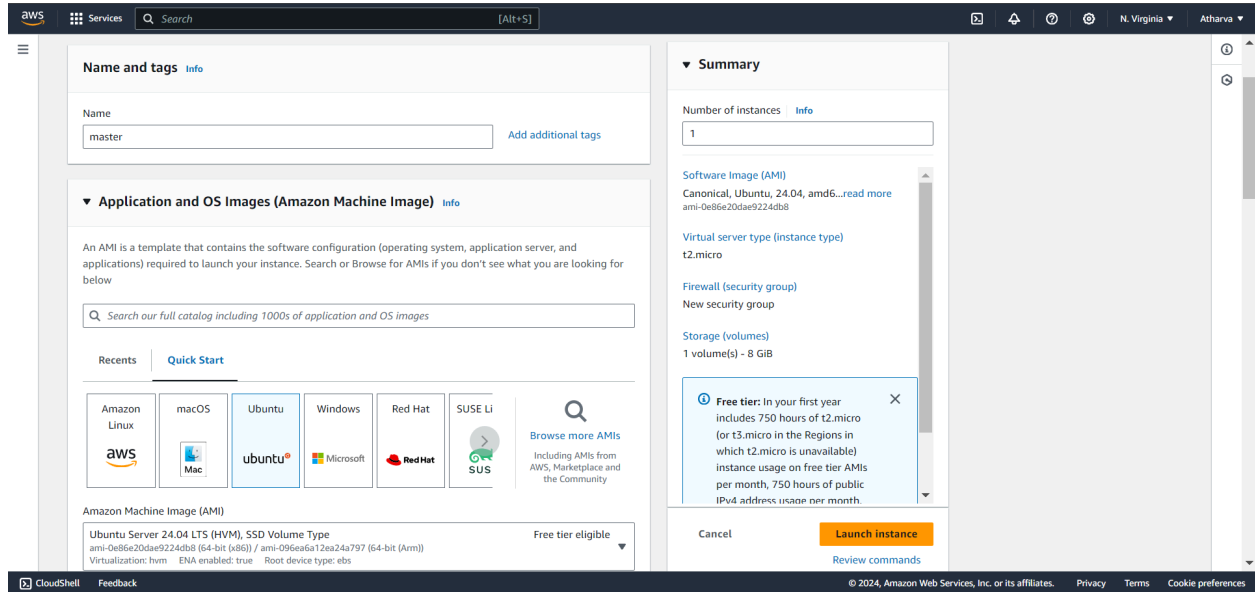
ATHARVA PRABHU D15A 43

Experiment 3

AIM: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Step 1: Prerequisites

1.1 Create 3 EC2 instances, one for the master node and two for the worker nodes.



1.2 Proceed with the following settings and create a new key pair as follows (use the same key pair for all the three nodes)

aws

Services

Search

[Alt+S]

Mumbai

Atharva

Name and tags

Info

Name

master

Add additional tags

▼ Application and OS Images (Amazon Machine Image)

Info

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Search our full catalog including 1000s of application and OS images

Quick Start

Amazon Linux

macOS

Ubuntu

Windows

Red Hat

SUSE Linux

Browse more AMIs

Amazon Machine Image (AMI)

Ubuntu Server 24.04 LTS (HVM), SSD Volume Type

ami-0522ab6e1ddc7055 (64-bit (x86)) / ami-0000791ba666add5 (64-bit (Arm))

Virtualization: hvm ENA enabled: true Root device type: ebs

Free tier eligible

▼ Summary

Number of instances

Info

1

Software Image (AMI)

Canonical, Ubuntu, 24.04, amd64...read more

ami-0522ab6e1ddc7055

Virtual server type (instance type)

t2.micro

Firewall (security group)

New security group

Storage (volumes)

1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 750 hours of public IPv4 address usage per month

Cancel

Launch instance

Create key pair

Key pair name

Key pairs allow you to connect to your instance securely.

three-tier-app

The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.

Key pair type

☒ RSA

RSA encrypted private and public key pair

☐ ED25519

ED25519 encrypted private and public key pair

Private key file format

☒ .pem

For use with OpenSSH

☐ .ppk

For use with PuTTY

⚠ When prompted, store the private key in a secure and accessible location on your computer. You will need it later to connect to your instance. [Learn more](#)

Cancel

Create key pair

aws Services Search (Alt+S)

EC2 > Instances > i-078492e2743594f43 > Connect to instance

Connect to instance [Info](#)

Connect to your instance i-078492e2743594f43 (master) using any of these options

EC2 Instance Connect

Session Manager

SSH client

EC2 serial console

Instance ID
i-078492e2743594f43 (master)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is three-tier-app.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
chmod 400 "three-tier-app.pem"
4. Connect to your instance using its Public DNS:
ec2-65-1-65-18.ap-south-1.compute.amazonaws.com

Command copied

```
ssh -i "three-tier-app.pem" ubuntu@ec2-65-1-65-18.ap-south-1.compute.amazonaws.com
```

Note: In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Cancel

1.3 Add port 6443 in each security group

Inbound rules [Info](#)

Security group rule ID	Type Info	Protocol Info	Port range Info	Source Info	Description - optional Info	
sgr-0cf20b6a9f8501fc6	Custom TCP	TCP	6443	Custom	<div>Q</div> <div>0.0.0.0/0 X</div> <div></div>	<div>Delete</div>
sgr-0e02c88e6fce1b710	SSH	TCP	22	Custom	<div>Q</div> <div>0.0.0.0/0 X</div> <div></div>	<div>Delete</div>

Add rule

1.4 After the instances have been created, copy the text given in the example part of each of the three instances into git bash.

```
C:\Users\Atharva\Downloads>ssh -i "three-tier-app.pem" ubuntu@ec2-65-1-65-18.ap-south-1.compute.amazonaws.com
Welcome to Ubuntu 24.04 LTS (GNU/Linux 6.8.0-1012-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/pro

System information as of Sat Sep 21 10:50:19 UTC 2024

System load:  0.11           Processes:            115
Usage of /:   22.9% of 6.71GB Users logged in:          0
Memory usage: 5%            IPv4 address for enx0: 172.31.46.220
Swap usage:   0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

Last login: Sat Sep 21 10:44:19 2024 from 49.36.97.186
To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.
```

Step 2: Run the following commands on both the master and worker nodes to prepare them for kubeadm.

```
# disable swap
sudo swapoff -a
```

```
# Create the .conf file to load the modules at bootup
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
overlay
br_netfilter
EOF
```

```
sudo modprobe overlay
sudo modprobe br_netfilter
```

```
# sysctl params required by setup, params persist across reboots
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
EOF
```

```
# Apply sysctl params without reboot
sudo sysctl --system
```

Install CRIO Runtime

```
sudo apt-get update -y
```

```
sudo apt-get install -y software-properties-common curl apt-transport-https  
ca-certificates gpg
```

```
sudo curl -fsSL https://pkgs.k8s.io/addons:/cri-o:/prerelease:/main/deb/Release.key |  
sudo gpg --dearmor -o /etc/apt/keyrings/cri-o-apt-keyring.gpg  
echo "deb [signed-by=/etc/apt/keyrings/cri-o-apt-keyring.gpg]  
https://pkgs.k8s.io/addons:/cri-o:/prerelease:/main/deb/ /" | sudo tee  
/etc/apt/sources.list.d/cri-o.list
```

```
sudo apt-get update -y
```

```
sudo apt-get install -y cri-o
```

```
sudo systemctl daemon-reload
```

```
sudo systemctl enable crio --now
```

```
sudo systemctl start crio.service
```

```
echo "CRI runtime installed successfully"
```

Add Kubernetes APT repository and install required packages

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.29/deb/Release.key | sudo gpg --dearmor  
-o /etc/apt/keyrings/kubernetes-apt-keyring.gpg  
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]  
https://pkgs.k8s.io/core:/stable:/v1.29/deb/ /' | sudo tee  
/etc/apt/sources.list.d/kubernetes.list
```

```
sudo apt-get update -y
```

```
sudo apt-get install -y kubelet="1.29.0-*" kubectl="1.29.0-*" kubeadm="1.29.0-*"
```

```
sudo apt-get update -y
```

```
sudo apt-get install -y jq
```

```
sudo systemctl enable --now kubelet
```

```
sudo systemctl start kubelet
```

```

ubuntu@ip-172-31-46-228:~$ # disable swap
sudo swapoff -a

# Create the .conf file to load the modules at bootup
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
overlay
br_netfilter
EOF

sudo modprobe overlay
sudo modprobe br_netfilter

# sysctl params required by setup, params persist across reboots
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
EOF

# Apply sysctl params without reboot
sudo sysctl --system

## Install CRI-O Runtime
sudo apt-get update -y
sudo apt-get install -y software-properties-common curl apt-transport-https ca-certificates gpg

sudo curl -fsSL https://pkgs.k8s.io/addons:/cri-o:/prerelease:/main/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/cri-o-apt-keyring.gpg
echo "deb [signed-by=/etc/apt/keyrings/cri-o-apt-keyring.gpg] https://pkgs.k8s.io/addons:/cri-o:/prerelease:/main/deb/ /" | sudo
re:/stable/v1.29/deb/ /' | sud
overlay
br_netfilter
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
* Applying /usr/lib/sysctl.d/10-apparmor.conf ...
* Applying /etc/sysctl.d/10-console-messages.conf ...
* Applying /etc/sysctl.d/10-ipv6-privacy.conf ...
* Applying /etc/sysctl.d/10-kernel-hardening.conf ...
* Applying /etc/sysctl.d/10-magic-sysrq.conf ...
* Applying /etc/sysctl.d/10-map-count.conf ...
* Applying /etc/sysctl.d/10-network-security.conf ...
* Applying /etc/sysctl.d/10-ptrace.conf ...
* Applying /etc/sysctl.d/10-zero-page.conf ...
* Applying /etc/sysctl.d/50-cloudimg-settings.conf ...
* Applying /usr/lib/sysctl.d/50-pid-max.conf ...
* Applying /etc/sysctl.d/99-cloudimg-ipv6.conf ...
* Applying /usr/lib/sysctl.d/99-protect-links.conf ...
* Applying /etc/sysctl.d/99-sysctl.conf ...

```

Step3: Run the above command only on master node

sudo kubeadm config images pull

sudo kubeadm init

mkdir -p "\$HOME"/.kube

sudo cp -i /etc/kubernetes/admin.conf "\$HOME"/.kube/config

sudo chown "\$(id -u)": "\$(id -g)" "\$HOME"/.kube/config

Network Plugin = calico

kubectl apply -f

<https://raw.githubusercontent.com/projectcalico/calico/v3.26.0/manifests/calico.yaml>

kubeadm token create --print-join-command

```

ubuntu@ip-172-31-46-228:~$ sudo kubeadm config images pull

sudo kubeadm init

mkdir -p "$HOME"/.kube
sudo cp -i /etc/kubernetes/admin.conf "$HOME"/.kube/config
sudo chown "$(id -u)": "$(id -g)" "$HOME"/.kube/config

# Network Plugin = calico
kubectctl apply -f https://raw.githubusercontent.com/projectcalico/calico/v3.26.0/manifests/calico.yaml

kubeadm token create --print-join-command
I0921 11:12:21.776399    3963 version.go:256] remote version is much newer: v1.31.0; falling back to: stable-1.29
[config/images] Pulled registry.k8s.io/kube-apiserver:v1.29.9
[config/images] Pulled registry.k8s.io/kube-controller-manager:v1.29.9
[config/images] Pulled registry.k8s.io/kube-scheduler:v1.29.9
[config/images] Pulled registry.k8s.io/kube-proxy:v1.29.9
[config/images] Pulled registry.k8s.io/coredns/coredns:v1.11.1
[config/images] Pulled registry.k8s.io/pause:3.9
[config/images] Pulled registry.k8s.io/etcd:3.5.10-0
I0921 11:12:40.995686    4304 version.go:256] remote version is much newer: v1.31.0; falling back to: stable-1.29
[init] Using Kubernetes version: v1.29.9
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
W0921 11:12:41.763411    4304 checks.go:835] detected that the sandbox image 'registry.k8s.io/pause:3.10' of the container runtime is inconsistent with that used by kubeadm. It is recommended that using "regist
ry.k8s.io/pause:3.9" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-46-228 kubernetes kubernetes.default kubernetes.default.svc kubernetes.default.svc.cluster.local] and IPs [10.96.0.1 172.31.46.228]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-46-228 localhost] and IPs [172.31.46.220 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-46-228 localhost] and IPs [172.31.46.220 127.0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key
[kubeconfig] Using kubeconfig folder "/etc/kubernetes"
[kubeconfig] Writing "admin.conf" kubeconfig file
[kubeconfig] Writing "super-admin.conf" kubeconfig file
[kubeconfig] Writing "kubelet.conf" kubeconfig file
[kubeconfig] Writing "controller-manager.conf" kubeconfig file
[kubeconfig] Writing "scheduler.conf" kubeconfig file
[etcd] Creating static Pod manifest for local etcd in "/etc/kubernetes/manifests"

```

You will get kubeadm token, Copy it.

Step 4: Run the above command only on worker nodes

sudo kubeadm reset pre-flight checks

sudo your-token --v=5

```

ubuntu@ip-172-31-46-212:~$ sudo kubeadm reset pre-flight checks
W0921 11:14:17.712669    3933 preflight.go:56] [reset] WARNING: Changes made to this host by 'kubeadm init' or 'kubeadm join' will be reverted.
[reset] Are you sure you want to proceed? [y/N]: yes
[preflight] Running pre-flight checks
W0921 11:14:20.535200    3933 removeetcdmember.go:106] [reset] No kubeadm config, using etcd pod spec to get data directory
[reset] Deleted contents of the etcd data directory: /var/lib/etcd
[reset] Stopping the kubelet service
[reset] Unmounting mounted directories in "/var/lib/kubelet"
[reset] Deleting contents of directories: [/etc/kubernetes/manifests /var/lib/kubelet /etc/kubernetes/pki]
[reset] Deleting files: [/etc/kubernetes/admin.conf /etc/kubernetes/super-admin.conf /etc/kubernetes/kubelet.conf /etc/kubernetes/bootstrap-kubelet.conf /etc/kubernetes/controller-manager.conf /etc/kubernetes/scheduler.conf]

The reset process does not clean CNI configuration. To do so, you must remove /etc/cni/net.d

The reset process does not reset or clean up iptables rules or IPVS tables.
If you wish to reset iptables, you must do so manually by using the "iptables" command.

If your cluster was setup to utilize IPVS, run ipvsadm --clear (or similar)
to reset your system's IPVS tables.

The reset process does not clean your kubeconfig files and you must remove them manually.

```

```

ubuntu@ip-172-31-36-212:~$ sudo kubeadm join 172.31.46.220:6443 --token k4psyh.nslg1yet9he59kd4 --discovery-token-ca-cert-hash sha256:88e7e9abf8f31f0333a9d2f7a688d6dbdf9612679ae45cf71bada8de069d4a292e --v=5
I0921 11:28:31.863878 4897 join.go:413] [preFlight] Found NodeName empty; using OS hostname as NodeName
I0921 11:28:31.864885 4897 initconfiguration.go:122] detected and using CRI socket: unix:///var/run/crio/crio.sock
[preFlight] Running pre-flight checks
I0921 11:28:31.864143 4897 preflight.go:93] [preFlight] Running general checks
I0921 11:28:31.864183 4897 checks.go:280] validating the existence of file /etc/kubernetes/kubelet.conf
I0921 11:28:31.864287 4897 checks.go:280] validating the existence of file /etc/kubernetes/bootstrap-kubelet.conf
I0921 11:28:31.864219 4897 checks.go:184] validating the container runtime
I0921 11:28:31.889669 4897 checks.go:639] validating whether swap is enabled or not
I0921 11:28:31.889763 4897 checks.go:370] validating the presence of executable crictl
I0921 11:28:31.889799 4897 checks.go:370] validating the presence of executable contrack
I0921 11:28:31.889819 4897 checks.go:370] validating the presence of executable ip
I0921 11:28:31.889843 4897 checks.go:370] validating the presence of executable iptables
I0921 11:28:31.889870 4897 checks.go:370] validating the presence of executable mount
I0921 11:28:31.889897 4897 checks.go:370] validating the presence of executable nsenter
I0921 11:28:31.889919 4897 checks.go:370] validating the presence of executable ebtables
I0921 11:28:31.889954 4897 checks.go:370] validating the presence of executable ethtool
I0921 11:28:31.889977 4897 checks.go:370] validating the presence of executable socat
I0921 11:28:31.889996 4897 checks.go:370] validating the presence of executable tc
I0921 11:28:31.890011 4897 checks.go:370] validating the presence of executable touch
I0921 11:28:31.890035 4897 checks.go:516] running all checks
I0921 11:28:31.103935 4897 checks.go:481] checking whether the given node name is valid and reachable using net.LookupHost
I0921 11:28:31.105638 4897 checks.go:685] validating kubelet version
I0921 11:28:31.162593 4897 checks.go:130] validating if the "kubelet" service is enabled and active
I0921 11:28:31.176512 4897 checks.go:283] validating availability of port 10250
I0921 11:28:31.176737 4897 checks.go:280] validating the existence of file /etc/kubernetes/pki/ca.crt
I0921 11:28:31.176765 4897 checks.go:430] validating if the connectivity type is via proxy or direct
I0921 11:28:31.176883 4897 checks.go:329] validating the contents of file /proc/sys/net/bridge/bridge-nf-call-iptables
I0921 11:28:31.176849 4897 checks.go:329] validating the contents of file /proc/sys/net/ipv4/ip_forward
I0921 11:28:31.176883 4897 join.go:532] [preFlight] Discovering cluster-info
I0921 11:28:31.176917 4897 token.go:180] [discovery] Created cluster-info discovery client, requesting info from "172.31.46.220:6443"
I0921 11:28:31.187676 4897 token.go:118] [discovery] Requesting info from "172.31.46.220:6443" again to validate TLS against the pinned public key
I0921 11:28:31.194531 4897 token.go:135] [discovery] Cluster info signature and contents are valid and TLS certificate validates against pinned roots, will use API Server "172.31.46.220:6443"
I0921 11:28:31.194680 4897 discovery.go:52] [discovery] Using provided TLSBootstrapToken as authentication credentials for the join process
I0921 11:28:31.194622 4897 join.go:846] [preFlight] Fetching init configuration
I0921 11:28:31.194629 4897 join.go:592] [preFlight] Retrieving KubeConfig objects
[preFlight] Reading configuration from the cluster...
[preFlight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
I0921 11:28:31.201989 4897 kubeproxy.go:55] attempting to download the KubeProxyConfiguration from ConfigMap "kube-proxy"
I0921 11:28:31.205146 4897 kubelet.go:74] attempting to download the KubeletConfiguration from ConfigMap "kubelet-config"
I0921 11:28:31.208370 4897 initconfiguration.go:114] skip CRI socket detection, fill with the default CRI socket unix:///var/run/containerd/containerd.sock
I0921 11:28:31.208595 4897 interface.go:432] Looking for default routes with IPv4 addresses
I0921 11:28:31.208617 4897 interface.go:437] Default route transits interface "enx0"
I0921 11:28:31.208751 4897 interface.go:289] Interface enx0 is up
I0921 11:28:31.208803 4897 interface.go:257] Interface "enx0" has 2 addresses :[172.31.36.212/20 fe80::75:41ff:fea5:afb1/64].
I0921 11:28:31.208819 4897 interface.go:224] Checking addr 172.31.36.212/20.
I0921 11:28:31.208829 4897 interface.go:231] IP found 172.31.36.212
I0921 11:28:31.208840 4897 interface.go:263] Found valid IPv4 address 172.31.36.212 for interface "enx0".
I0921 11:28:31.208849 4897 interface.go:440] Found active IP 172.31.36.212
I0921 11:28:31.215092 4897 preflight.go:184] [preFlight] Running configuration dependant checks
I0921 11:28:31.215028 4897 controlplaneprepare.go:225] [download-certs] Skipping certs download

```

Step5: Run the given command to verify cluster creation

kubectl get nodes

```

ubuntu@ip-172-31-46-220:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE    VERSION
ip-172-31-36-212                    Ready    <none>    47s    v1.29.0
ip-172-31-46-220                    Ready    control-plane  16m    v1.29.0
ip-172-31-47-26                     Ready    <none>    29s    v1.29.0

```